

2024 Spring Semester

Operating System Final Exam (40pts)

Deadlock (10pts)

1. Assume that a computer system currently runs five processes which are in the running state. Each process is single-threaded and provided its maximum resource requirements in its creation. Below is a resource allocation table taken for the system:

Process ID	Allocation	Maximum	Available
	A B C	A B C	A B C
P1	1 1 2	3 3 3	2 1 2
P2	1 2 2	3 2 5	
P3	1 2 1	1 3 1	
P4	0 1 0	3 1 5	
P5	0 0 0	4 3 0	

a. Describe validity of the processes first, and then explain which state the current system is by using Banker's algorithm. (3pts)

-> First of all, Banker's algorithm detects the unsafe state, not the deadlock state.

The very first job to do is checking whether the processes requests lower or equal to the maximum system resource instances. By adding the allocated and available resource instances, the system has 5, 7, and 7 instances for A, B, and C resources. All processes' maximum instances are less than them, so the processes seem valid.

Second, by applying the Banker's algorithm, let's consider index starts from 1 in each iteration. Then, the sequence for safety check exists as <P3 (2,2,2), P1 (3,3,4), P2 (4,5,6), P4 (5,7,7), P5 (5,7,7)>, which means the system is in the safe state.

b. If P5 requests two resource instances of C, explain how to handle this request as if you were an OS of this system. Also, explain whether granting this request would make the system be deadlocked, too. (3pts)

-> P5 claimed no resource instances at its creation, therefore this request is definitely invalid so it should not be granted.

However, if we apply the deadlock detection algorithm for the single request of 0, 0, 2 by P5, the sequence of <P3, P1, P2, P4, P5> still does not change, therefore the system becomes not deadlocked though the request is invalid.

c. In the deadlock prevention, OS implements its structure to inherently block one of four deadlock requirements. Please list up the four requirements and explain how to block each

requirement with its side effect in the general OS. (4pts)

-> Mutual Exclusion: making all OS components be sharable could violate this requirement, but some components in the general OS, e.g. mutex lock, are intrinsically nonsharable, which requires mutual exclusion.

Hold and Wait: executing processes if and only if all required resource instances are allocated. This method definitely blocks the deadlock, however, resource utilization becomes low and starvation possibly occurs.

No Preemption: Enabling a process steals resource instances from the other process holding the required resources. Same to the Hold and Wait case, starvation can occur if the victim selection is not adequately designed.

Circular Wait: Numbering all resources and granting resource requests in one direction only, not to make any cycles in resource requests. This restricts all processes' behavior to acquire resources in the fixed order, therefore reduces overall efficiency.

2. Main Memory Management (10pts)

a. Please answer with **True or False** if you know the answer. (5pts, **right answer 1pt, blank 0pt, wrong answer -1pt**) (3pts)

(i) If a process repeatedly calls for the system libraries which are dynamically linked, then the stub finds and runs the relevant codes.

-> False, the first call lets the stub find the relevant address, but the indirection table is updated during this first call and subsequent calls do not require the stub.

(ii) Paging is a memory management technique that divides both physical memory and logical address space of a process into fixed-size blocks.

-> True

(iii) By using the paging with segmentation, OS can run a process that exceeds the physical main memory size.

-> False. Virtual memory enables this functionality.

(iv) Demand paging requires a whole process image to be uploaded in the main memory for running.

-> False. It just uploads pages which is referenced by a process.

(v) Increasing the page size also increases the page table size.

-> False. Increasing the page size for a system decreases the number of page table entry so the page table size decreases.

b. Assume that a system implemented a three-level hierarchical paging system. Here, the


system needs 100 ns for main memory access while the cache access time is 10 ns. If the three-level hierarchical paging system is implemented with 80% hit ratio on average, answer its effective access time with calculation steps. (2pt)

-> For the cache hit, the memory access consists of $m + e$, 110 ns.

For the cache miss, the initial cache access, three page table accesses, and the main memory access are required as $4m + e$, 410 ns.

Therefore, overall EAT becomes $(m+e) * 0.8 + (4m+e) * 0.2 = 110\text{ns} * 0.8 + 410 * 0.2 = 88\text{ns} + 82\text{ns} = 170\text{ns}$

c. Assume that you are going to implement a hashed page table system. To make it more efficient than the three-level hierarchical paging system in 2-b, you have designed a hash function that minimizes the hash collisions so your hash function makes only two collisions in each entry while the hash function calculation takes 5 ns. If the first chain hits 80% on average, what should be the minimum cache hit ratio to outperform the 3-level paging system in 2-b? Answer with EAT calculation. (3pts)



-> Similarly, the cache hit requires one memory access and one cache access, $m + e$.

For the cache miss, two scenarios exist:

1) 1st chain hit: the initial cache access occurs, hash function calculation, hash table access, first chain comparison, main memory access

$$e + h + 3m$$

2) 2nd chain hit: the initial cache access occurs, hash function calculation, hash table access, first chain comparison, second chain comparison, main memory access

$$e + h + 4m$$

$$\text{EAT} = (m + e)a + (e + h + 3m)(1-a)0.8 + (e + h + 4m)(1-a)0.2 \leq 170$$

$$= (m - h - 3.2m)a + (e + h + 3.2m)$$

$$= (-2.2m - h)a + (e + h + 3.2m) \leq 170$$

$$= -225a + 335 \leq 170$$

$$= -225a \leq -155$$

$$a \geq 155/225 = 31/45 \approx 68.89\%$$

[illegible]

-> $55 + 11 + 44 = 110 \text{ ms}$

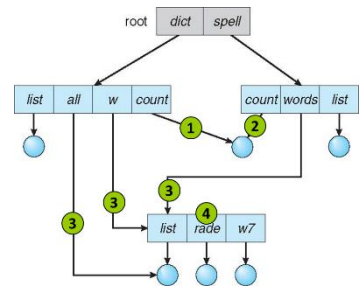
[illegible]

-> 8 page faults, $55 + 11 + 42 = 108$ ms

[illegible]

4. File-system (8pts)

a. In the right figure, let's assume that a user deleted (1) pointer. Then, answer the name of (2) and one solution not to make (2) in the file system. Also, if a user deleted (3) pointers, then explain what OS does for handling (4) and answer the procedure's name. (2pts)



-> (2) -> Dangling pointer

Solutions: by keeping backpointers using the daisy-chain structure, so if a user deletes a file then OS can delete all pointers for the file or delete the commanded link only.

By keeping a counter of pointers in a file or a directory, the deleting operation decreases the counter and the file is completely deleted if the counter becomes zero.

(4) OS traverses all files and directories first while marking them, and scan files and directories that cannot be accessed from its root directory. This procedure is called as garbage collection.

b. Let's say a user process runs two following codes:

```
myFile = fopen("c:\OS\finalExam.txt", "r");
```

```
fscanf(myFile, "%s\n", inString);
```

Please explain the procedure of access, open, and read the file using the in-memory file system structures. (4pt)

-> At the secondary drive mounting, the directory structure is uploaded in the kernel memory. At the first code, the function call of `fopen()` is delivered to the system call `open()` via dual-mode operation, and the kernel thread resolves the path to find where the file control block of "finalExam.txt" exists. After uploading the file control block to the kernel space from the disk through I/O, the kernel stores the pointer to the file control block in the system-wide open file table and updates the process-wide open-file table. After returning to the user mode, the process receives the index of the process's open-file table as a file handle.

Then, for `fscanf`, a user process sends the file handle with parameters. The relevant system call, `read()`, runs after changed to the kernel mode and refers the process-wide and system-wide open-file tables using the file handle to find the address of the file control block. With the parameters, the routine of `read()` now calculates the relative disk location of the target data and sends it to the disk I/O with the address of `inString` and change the process state as waiting. After the reading I/O, the process is woken up and return to the user mode, and the user process can access the data through `inString` variable.

c. In comparison to the contiguous allocation, explain the advantages of the extent-based systems. Also, describe what happens in the extent-based systems for the long-time usage. (2pts)

-> Extent-based systems divide the contiguous file into several sub-contiguous parts so mostly provide an easy way of being allocated in the secondary disk with no external fragmentation while keeping fast read and write speeds coming from the contiguous allocation. However, the

long-time usage leads to the increased number of extents which requires more data to store the metadata of extents and potentially slower read and write speeds to move the disk magnetic arm to the different extent locations.

Good Luck!